

Baseline V2 Final Review and Split Report

Final review counts

Status	Count
approved	412
rejected	118
relabel	0
pending	0

Class viability

Locked rule: at least 10 approved photos and at least 8 independent approved perceptual groups.

Class	Total	Approved	Rejected	Relabeled	Groups	Viable	Reason
filling_station	21	21	0	0	21	yes	—
gap_in_median	13	13	0	0	13	yes	—
give_way	21	18	3	0	16	yes	—
hairpin_bend_ahead	23	23	0	0	19	yes	—
left_curve_ahead	29	26	3	0	22	yes	—
major_road_ahead	16	16	0	0	15	yes	—
maximum_speed_limit_30_km_h	19	16	3	0	15	yes	—
maximum_speed_limit_40_km_h	23	20	3	0	19	yes	—
maximum_speed_limit_80_km_h	19	19	0	0	17	yes	—
no_entry	17	14	3	0	12	yes	—
no_right_turn	11	11	0	0	11	yes	—
one_way	18	0	18	0	0	no	approved photographs 0 < 10; approved perceptual groups 0 < 8
pass_either_side	12	11	1	0	11	yes	—
pedestrian_crossing_ahead	80	63	17	0	56	yes	—

Class	Total	Approved	Rejected	Relabeled	Groups	Viable	Reason
right_curve_ahead	33	31	2	0	29	yes	—
road_hump	61	53	8	0	48	yes	—
school_ahead	43	38	5	0	35	yes	—
side_road_right	19	19	0	0	17	yes	—
stop	36	0	36	0	0	no	approved photographs 0 < 10; approved perceptual groups 0 < 8
y_junction	16	0	16	0	0	no	approved photographs 0 < 10; approved perceptual groups 0 < 8

Group-safe split

Per class, sort groups deterministically, apply a seed-42 stable shuffle for tie ordering, then use dynamic programming to assign each complete perceptual group to train/validation/test. Select the reachable non-empty allocation minimizing squared sample-count deviation from 70/15/15; deterministic lexical assignment breaks exact ties.

Split	Samples	Groups	Target	Actual	Deviation
train	287	251	70.00%	69.66%	-0.34%
validation	62	62	15.00%	15.05%	+0.05%
test	63	63	15.00%	15.29%	+0.29%

Per-class split distribution

Class	Train	Validation	Test	Total
filling_station	15	3	3	21
gap_in_median	9	2	2	13
give_way	12	3	3	18
hairpin_bend_ahead	16	3	4	23
left_curve_ahead	18	4	4	26
major_road_ahead	11	2	3	16
maximum_speed_limit_30_kmh	11	2	3	16
maximum_speed_limit_40_kmh	14	3	3	20
maximum_speed_limit_80_kmh	13	3	3	19
no_entry	10	2	2	14
no_right_turn	8	2	1	11
pass_either_side	8	2	1	11
pedestrian_crossing_ahead	44	9	10	63
right_curve_ahead	22	5	4	31
road_hump	37	8	8	53
school_ahead	26	6	6	38
side_road_right	13	3	3	19

Perceptual-group leakage checks

Check	Result	Details
duplicate_review_ids	PASS	0
duplicate_source_image_ids	PASS	0
exact_duplicate_groups	PASS	0
near_duplicate_pair_group_inconsistencies	PASS	0
cross_label_perceptual_group_conflicts	PASS	0
split_perceptual_group_overlap	PASS	0
split_source_image_overlap	PASS	0
split_review_id_overlap	PASS	0

Warnings and limitations

- no_right_turn and pass_either_side each have one test sample; their per-class test metrics will be unstable.
- Approved class sizes range from 11 to 63; calculate any class weights from the training split only.
- Dataset B was observed during Baseline V1 and is not a pristine unseen V2 benchmark.
- The test split is held out and must not be used for model selection; use validation macro F1.
- Older remaining-review exports are interim snapshots; final_review_summary.json and v2_split_summary.json are authoritative.

No Baseline V2 model was trained or evaluated while creating this bundle.